


DEADLOCKS

Deadlock is a situation where a set of processes are blocked because each process is

- holding a resource and
- waiting for another resource held by some other process.

System Model

- ✓ A system consist of finite number of resources. (For ex: memory, printers, CPUs).
- ✓ These resources are distributed among number of processes.
- ✓ A process must
 - request a resource before using it and
 - release the resource after using it.
- ✓ The process can request any number of resources to carry out a given task.
- ✓ The total number of resource requested must not exceed the total number of resources available.
- ✓ In normal operation, a process must perform following tasks in sequence:
 - **Request**
If the request cannot be granted immediately (for ex: the resource is being used by another process), then the requesting-process must wait for acquiring the resource.
For example: open(), malloc(), new(), and request()
 - **Use**
The process uses the resource.
For example: prints to the printer or reads from the file.
 - **Release**
The process releases the resource.
So that, the resource becomes available for other processes.
For example: close(), free(), delete(), and release().
- ✓ A set of processes is deadlocked when every process in the set is waiting for a resource that is currently allocated to another process in the set.
- ✓ Deadlock may involve different types of resources.
- ✓ As shown in figure below,



- ✓ Both processes P1 & P2 need resources to continue execution.
- ✓ P1 requires additional resource R1 and is in possession of resource R2.
- ✓ P2 requires additional resource R2 and is in possession of R1.
Thus, neither process can continue.
- ✓ Multithread programs are good candidates for deadlock because in figure they compete for shared resources.

Deadlock Characterization

In a deadlock, processes never finish executing, and system resources are tied up, preventing other jobs from starting.

Necessary Conditions

There are four conditions that are necessary to achieve deadlock:

1. Mutual Exclusion

At least one resource must be held in a non-sharable mode.

If any other process requests this resource, then the requesting-process must wait for the resource to be released.

2. Hold and Wait

A process must be simultaneously

→ holding at least one resource and

→ waiting to acquire additional resources held by the other process.

3. No Preemption

Once a process is holding a resource (i.e. once its request has been granted), then that resource cannot be taken away from that process until the process voluntarily releases it.

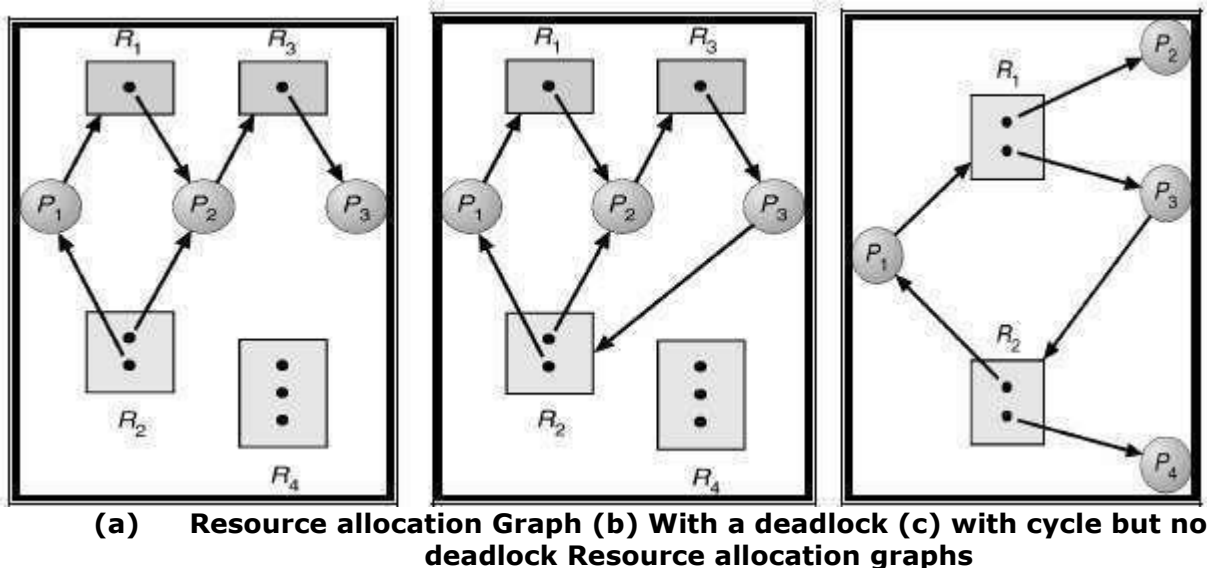
4. Circular Wait

A set of processes $\{ P_0, P_1, P_2, \dots, P_N \}$ must exist such that P_0 is waiting for a resource that is held by P_1

P_1 is waiting for a resource that is held by P_2 , and so on

Resource-Allocation-Graph

- ✓ The resource-allocation-graph (RAG) is a directed graph that can be used to describe the deadlock situation.
- ✓ RAG consists of a
 - set of vertices (V) and
 - set of edges (E).
- ✓ V is divided into two types of nodes
 - $P = \{P_1, P_2, \dots, P_n\}$ i.e., set consisting of all active processes in the system.
 - $R = \{R_1, R_2, \dots, R_n\}$ i.e., set consisting of all resource types in the system.
- ✓ E is divided into two types of edges:
 - **Request Edge**
A directed-edge $P_i \rightarrow R_j$ is called a request edge.
 $P_i \rightarrow R_j$ indicates that process P_i has requested a resource R_j .
 - **Assignment Edge**
A directed-edge $R_j \rightarrow P_i$ is called an assignment edge.
 $R_j \rightarrow P_i$ indicates that a resource R_j has been allocated to process P_i .
- ✓ Suppose that process P_i requests resource R_j .
Here, the request for R_j from P_i can be granted only if the converting request-edge to assignment-edge do not form a cycle in the resource-allocation graph.
- ✓ Pictorially,
 - We represent each process P_i as a circle.
 - We represent each resource-type R_j as a rectangle.
- ✓ As shown in below figures, the RAG illustrates the following 3 situation :
 - RAG with a deadlock
 - RAG with a cycle and deadlock
 - RAG with a cycle but no deadlock



Methods for Handling Deadlocks

- ✓ There are three ways of handling deadlocks:
 - Deadlock prevention or avoidance - Do not allow the system to get into a deadlocked state.
 - Deadlock detection and recovery - Abort a process or preempt some resources when deadlocks are detected.
 - Ignore the problem all together - If deadlocks only occur once a year or so, it may be better to simply let them happen and reboot the system.
- ✓ In order to avoid deadlocks, the system must have additional information about all processes.
- ✓ In particular, the system must know what resources a process will or may request in the future.
- ✓ Deadlock detection is fairly straightforward, but deadlock recovery requires either aborting processes
- ✓ or preempting resources.
- ✓ If deadlocks are neither prevented nor detected, then when a deadlock occurs the system will gradually slow down.

Deadlock-Prevention

Deadlocks can be eliminated by preventing at least one of the four required conditions:

1. Mutual exclusion
2. Hold-and-wait
3. No preemption
4. Circular-wait.

Mutual Exclusion

- ✓ This condition must hold for non-sharable resources.
- ✓ For example:
A printer cannot be simultaneously shared by several processes.
- ✓ On the other hand, shared resources do not lead to deadlocks.
- ✓ For example:
Simultaneous access can be granted for read-only file.
- ✓ A process never waits for accessing a sharable resource.

- ✓ In general, we cannot prevent deadlocks by denying the mutual-exclusion condition because some resources are non-sharable by default.

Hold and Wait

- ✓ To prevent this condition:
The processes must be prevented from holding one or more resources while simultaneously waiting for one or more other resources.
- ✓ There are several solutions to this problem.
- ✓ For example:
Consider a process that
 - copies the data from a tape drive to the disk
 - sorts the file and
 - then prints the results to a printer.
- ✓ **Protocol-1**
 - Each process must be allocated with all of its resources before it begins execution.
 - All the resources (tape drive, disk files and printer) are allocated to the process at the beginning.
- ✓ **Protocol-2**
 - A process must request a resource only when the process has none.
 - Initially, the process is allocated with tape drive and disk file.
 - The process performs the required operation and releases both tape drive and disk file.
 - Then, the process is again allocated with disk file and the printer
 - Again, the process performs the required operation & releases both disk file and the printer.
- ✓ **Disadvantages** of above 2 methods:
 - Resource utilization may be low, since resources may be allocated but unused for a long period.
 - Starvation is possible.

No Preemption

- ✓ To prevent this condition: the resources must be preempted.
- ✓ There are several solutions to this problem.
- ✓ **Protocol-1**
 - If a process is holding some resources and requests another resource that cannot be immediately allocated to it, then all resources currently being held are preempted.
 - The preempted resources are added to the list of resources for which the process is waiting.
 - The process will be restarted only when it regains the old resources and the new resources that it is requesting.
- ✓ **Protocol-2**
 - When a process request resources, we check whether they are available or not.
- ✓ These 2 protocols may be applicable for resources whose states are easily saved and restored, such
- ✓ as registers and memory.
- ✓ But, these 2 protocols are generally not applicable to other devices such as printers and tape drives.

```

If (resources are
available) then
{
    allocate resources to the process
}
else
{
    If (resources are allocated to waiting
    process) then
    {
        preempt the resources from the waiting
        process allocate the resources to the
        requesting-process the requesting-process
        must wait
    }
}

```

Circular-Wait

- ✓ Deadlock can be prevented by using the following 2 protocol:
- ✓ Protocol-1
 - Assign numbers all resources.
 - Require the processes to request resources only in increasing/decreasing order.
- ✓ Protocol-2
 - Require that whenever a process requests a resource, it has released resources with a lower number.
- ✓ One big challenge in this scheme is determining the relative ordering of the different resources.

Deadlock Avoidance


- ✓ The general idea behind deadlock avoidance is to prevent deadlocks from ever happening.
- ✓ Deadlock-avoidance algorithm
 - requires more information about each process, and
 - tends to lead to low device utilization.
- ✓ For example:
 - In simple algorithms, the scheduler only needs to know the maximum number of each resource that a process might potentially use.
 - In complex algorithms, the scheduler can also take advantage of the schedule of exactly what resources may be needed in what order.
- ✓ A deadlock-avoidance algorithm dynamically examines the resources allocation state to ensure that a circular-wait condition never exists.
- ✓ The resource-allocation state is defined by
 - the number of available and allocated resources and
 - the maximum demand of each process.

Safe State

- ✓ A state is safe if the system can allocate all resources requested by all processes without entering a deadlock state.
- ✓ A state is safe if there exists a safe sequence of processes {P0, P1, P2, ..., PN}

such that the requests of each process(P_i) can be satisfied by the currently available resources.


- ✓ If a safe sequence does not exist, then the system is in an unsafe state, which may lead to deadlock.
- ✓ All safe states are deadlock free, but not all unsafe states lead to deadlocks.



Safe, unsafe, and deadlock state spaces

Resource-Allocation-Graph Algorithm

- ✓ If resource categories have only single instances of their resources, then deadlock states can be detected by cycles in the resource-allocation graphs.
- ✓ In this case, unsafe states can be recognized and avoided by augmenting the resource-allocation graph with claim edges (denoted by a dashed line).
- ✓ Claim edge $P_i \rightarrow R_j$ indicated that process P_i may request resource R_j at some time in future.
- ✓ The important steps are as below:
 - When a process P_i requests a resource R_j , the claim edge $P_i \rightarrow R_j$ is converted to a request edge.
 - Similarly, when a resource R_j is released by the process P_i , the assignment edge $R_j \rightarrow P_i$ is reconverted as claim edge $P_i \rightarrow R_j$.
 - The request for R_j from P_i can be granted only if the converting request edge to assignment edge do not form a cycle in the resource allocation graph.
- ✓ To apply this algorithm, each process P_i must know all its claims before it starts executing.
- ✓ Conclusion:
 - If no cycle exists, then the allocation of the resource will leave the system in a safe state.
 - If cycle is found, system is put into unsafe state and may cause a deadlock.
- ✓ For example: Consider a resource allocation graph shown in Figure 3.5(a).
 - Suppose P_2 requests R_2 .
 - Though R_2 is currently free, we cannot allocate it to P_2 as this action will create a cycle in the graph as shown in Figure 3.5(b).
 - This cycle will indicate that the system is in unsafe state: because, if P_1 requests R_2 and P_2 requests R_1 later, a deadlock will occur.



Banker's Algorithm

- ✓ This algorithm is applicable to the system with multiple instances of each resource types.
- ✓ However, this algorithm is less efficient than the resource-allocation-graph algorithm.
- ✓ When a process starts up, it must declare the maximum number of resources that it may need.
- ✓ This number may not exceed the total number of resources in the system.
- ✓ When a request is made, the system determines whether granting the request would leave the system in a safe state.

If the system is in a safe state,

the resources are allocated;

else

the process must wait until some other process releases enough resources.

- ✓ Assumptions:

Let n = number of processes in the system

Let m = number of resource types.

- ✓ Following data structures are used to implement the banker's algorithm.

- Available [m]
 - This vector indicates the no. of available resources of each type.
 - If $\text{Available}[j]=k$, then k instances of resource type R_j is available.
- Max [n][m]
 - This matrix indicates the maximum demand of each process of each resource.
 - If $\text{Max}[i,j]=k$, then process P_i may request at most k instances of resource type R_j .
- Allocation [n][m]
 - This matrix indicates no. of resources currently allocated to each process.
 - If $\text{Allocation}[i,j]=k$, then P_i is currently allocated k instances of R_j .
- Need [n][m]
 - This matrix indicates the remaining resources need of each process.
 - If $\text{Need}[i,j]=k$, then P_i may need k more instances of resource R_j to complete its task.
 - So, $\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$

- ✓ The Banker's algorithm has two parts:
 - 1) Safety Algorithm
 - 2) Resource – Request Algorithm

Safety Algorithm

- ✓ This algorithm is used for finding out whether a system is in safe state or not.
- ✓ Assumptions:
 - Work is a working copy of the available resources, which will be modified during the analysis.
 - Finish is a vector of boolean values indicating whether a particular process can finish.

Step 1:

Let Work and Finish be two vectors of length m and n respectively.
 Initialize:
 Work = Available
 Finish[i] = false for $i=1,2,3,\dots,n$

Step 2:

Find an index(i) such that both
 a) Finish[i] = false
 b) Need i \leq Work.
 If no such i exist, then go to step 4

Step 3:

Set:
 Work = Work +
 Allocation(i) Finish[i]

Resource-Request Algorithm

- ✓ This algorithm determines if a new request is safe, and grants it only if it is safe to do so.
- ✓ When a request is made (that does not exceed currently available resources), pretend it has been granted, and then see if the resulting state is a safe one. If so, grant the request, and if not, deny the request.
- ✓ Let Request(i) be the request vector of process P_i .
- ✓ If Request(i)[j]=k, then process P_i wants K instances of the resource type R_j .

Step 1:

If $\text{Request}(i) \leq \text{Need}(i)$ then
go to step 2
else

raise an error condition, since the process has exceeded its maximum claim.

Step 2:

If $\text{Request}(i) \leq \text{Available}$ then
go to step 3
else

Pi must wait, since the resources are not available.

Step 3:

If the system want to allocate the requested resources to process Pi then modify the state as follows:

$\text{Available} = \text{Available} - \text{Request}(i)$

$\text{Allocation}(i) = \text{Allocation}(i) +$

$\text{Request}(i)$ $\text{Need}(i) = \text{Need}(i) -$

$\text{Request}(i)$

Step 4:

If the resulting resource-allocation state is safe, then i) transaction is complete and

ii) Pi is allocated its resources.

Step 5:

If the new state is unsafe,

then i) Pi must wait for $\text{Request}(i)$ and

ii) old resource-allocation state is restored.

Examples :

Question: Consider the following snapshot of a system:

	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	5	3	3	3	2
P1	2	0	0	3	2	2			
P2	3	0	3	9	0	2			
P3	2	1	1	2	2	2			
P4	0	0	2	4	3	3			

Answer the following questions using Banker's algorithm.

i) What is the content of the matrix need?

ii) Is the system in a safe state?

iii) If a request from process P1 arrives for (1 0 2) can the request be granted immediately?

Solution (i):

- The content of the matrix Need is given by
Need = Max - Allocation
- So, the content of Need Matrix is:

	Need		
	A	B	C
P0	7	4	3
P1	1	2	2
P2	6	0	0
P3	0	1	1
P4	4	3	1

Solution (ii):

- Applying the Safety algorithm on the given system,

Step 1: Initialization

Work = Available i.e. Work = 3 3 2

.....P0.....P1.....P2.....P3.....P4.....

Finish = | false | false | false | false | false |

Step 2: For i=0

Finish[P0] = false and Need[P0] ≤ Work i.e. (7 4 3) ≤ (3 3 2) → false
So P0 must wait.

Step 2: For i=1

Finish[P1] = false and Need[P1] ≤ Work i.e. (1 2 2) ≤ (3 3 2) → true
So P1 must be kept in safe sequence.

Step 3: Work = Work + Allocation[P1] = (3 3 2) + (2 0 0) = (5 3 2)

.....P0.....P1.....P2.....P3.....P4.....

Finish = | false | true | false | false | false |

Step 2: For i=2

Finish[P2] = false and Need[P2] ≤ Work i.e. (6 0 0) ≤ (5 3 2) → false
So P2 must wait.

Step 2: For i=3

Finish[P3] = false and Need[P3] ≤ Work i.e. (0 1 1) ≤ (5 3 2) → true
So P3 must be kept in safe sequence.

Step 3: Work = Work + Allocation[P3] = (5 3 2) + (2 1 1) = (7 4 3)

.....P0.....P1.....P2.....P3.....P4.....

Finish = | false | true | false | true | false |

Step 2: For $i=4$

$\text{Finish}[P4] = \text{false}$ and $\text{Need}[P4] \leq \text{Work}$ i.e. $(4 \ 3 \ 1) \leq (7 \ 4 \ 3) \rightarrow \text{true}$
So $P4$ must be kept in safe sequence.

Step 3: $\text{Work} = \text{Work} + \text{Allocation}[P4] = (7 \ 4 \ 3) + (0 \ 0 \ 2) = (7 \ 4 \ 5)$

.....P0.....P1.....P2.....P3.....P4.....
 $\text{Finish} = [\text{false} \ | \ \text{true} \ | \ \text{false} \ | \ \text{true} \ | \ \text{true} \]$

Step 2: For $i=0$

$\text{Finish}[P0] = \text{false}$ and $\text{Need}[P0] \leq \text{Work}$ i.e. $(7 \ 4 \ 3) \leq (7 \ 4 \ 5) \rightarrow \text{true}$
So $P0$ must be kept in safe sequence.

Step 3: $\text{Work} = \text{Work} + \text{Allocation}[P0] = (7 \ 4 \ 5) + (0 \ 1 \ 0) = (7 \ 5 \ 5)$

.....P0.....P1.....P2.....P3.....P4.....
 $\text{Finish} = [\text{true} \ | \ \text{true} \ | \ \text{false} \ | \ \text{true} \ | \ \text{true} \]$

Step 2: For $i=2$

$\text{Finish}[P2] = \text{false}$ and $\text{Need}[P2] \leq \text{Work}$ i.e. $(6 \ 0 \ 0) \leq (7 \ 5 \ 5) \rightarrow \text{true}$
So $P2$ must be kept in safe sequence.

Step 3: $\text{Work} = \text{Work} + \text{Allocation}[P2] = (7 \ 5 \ 5) + (3 \ 0 \ 2) = (10 \ 5 \ 7)$

.....P0.....P1.....P2.....P3.....P4.....
 $\text{Finish} = [\text{true} \ | \ \text{true} \ | \ \text{true} \ | \ \text{true} \ | \ \text{true} \]$

Step 4: $\text{Finish}[Pi] = \text{true}$ for $0 \leq i \leq 4$

Hence, the system is currently in a safe state.

The safe sequence is $\langle P1, P3, P4, P0, P2 \rangle$.

the system is currently in a safe state.

Solution (iii): $P1$ requests $(1 \ 0 \ 2)$ i.e. $\text{Request}[P1] = 1 \ 0 \ 2$

- To decide whether the request is granted, we use Resource Request algorithm.

Step 1: $\text{Request}[P1] \leq \text{Need}[P1]$ i.e. $(1 \ 0 \ 2) \leq (1 \ 2 \ 2) \rightarrow \text{true}$.

Step 2: $\text{Request}[P1] \leq \text{Available}$ i.e. $(1 \ 0 \ 2) \leq (3 \ 3 \ 2) \rightarrow \text{true}$.

Step 3: $\text{Available} = \text{Available} - \text{Request}[P1] = (3 \ 3 \ 2) - (1 \ 0 \ 2) = (2 \ 3 \ 0)$

$\text{Allocation}[P1] = \text{Allocation}[P1] + \text{Request}[P1] = (2 \ 0 \ 0) + (1 \ 0 \ 2) = (3 \ 0 \ 2)$

$\text{Need}[P1] = \text{Need}[P1] - \text{Request}[P1] = (1 \ 2 \ 2) - (1 \ 0 \ 2) = (0 \ 2 \ 0)$

- We arrive at the following new system state:

	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	5	3	2	3	0
P1	3	0	2	3	2	2			
P2	3	0	2	9	0	2			
P3	2	1	1	2	2	2			
P4	0	0	2	4	3	3			

- The content of the matrix
 $\text{Need} = \text{Max} -$
- So, the content of Need

Need is given by
Allocation
Matrix is:

	Need		
	A	B	C
P0	7	4	3
P1	0	2	0
P2	6	0	0
P3	0	1	1
P4	4	3	1

- To determine whether this new system state is safe, we again execute Safety algorithm.

Step 1: Initialization

Here, $m=3$, $n=5$

Work = Available i.e. Work = 2 3 0

.....P0.....P1.....P2.....P3.....P4.....

Finish = | false | false | false | false | false |

Step 2: For $i=0$

Finish[P0] = false and Need[P0] ≤ Work i.e. (7 4 3) ≤ (2 3 0) → false
So P0 must wait.

Step 2: For $i=1$

Finish[P1] = false and Need[P1] ≤ Work i.e. (0 2 0) ≤ (2 3 0) → true
So P1 must be kept in safe sequence.

Step 3: Work = Work + Allocation[P1] = (2 3 0) + (3 0 2) = (5 3 2)

.....P0.....P1.....P2.....P3.....P4.....

Finish = | false | true | false | false | false |

Step 2: For $i=2$

Finish[P2] = false and Need[P2] ≤ Work i.e. (6 0 0) ≤ (5 3 2) → false
So P2 must wait.

Step 2: For $i=3$

Finish[P3] = false and Need[P3] ≤ Work i.e. (0 1 1) ≤ (5 3 2) → true
So P3 must be kept in safe sequence.

Step 3: Work = Work + Allocation[P3] = (5 3 2) + (2 1 1) = (7 4 3)

.....P0.....P1.....P2.....P3.....P4.....

Finish = | false | true | false | true | false |

Step 2: For $i=4$

Finish[P4] = false and Need[P4] ≤ Work i.e. (4 3 1) ≤ (7 4 3) → true
So P4 must be kept in safe sequence.

Step 3: Work = Work + Allocation[P4] = (7 4 3) + (0 0 2) = (7 4 5)

.....P0.....P1.....P2.....P3.....P4.....

Finish = | false | true | false | true | true |

Step 2: For $i=0$

Finish[P0] = false and Need[P0] ≤ Work i.e. (7 4 3) ≤ (7 4 5) → true
So P0 must be kept in safe sequence.

Step 3: Work = Work + Allocation[P0] = (7 4 5) + (0 1 0) = (7 5 5)

.....P0.....P1.....P2.....P3.....P4.....

Finish = | true | true | false | true | true |

Step 2: For $i=2$

Finish[P2] = false and Need[P2] ≤ Work i.e. (6 0 0) ≤ (7 5 5) → true
So P2 must be kept in safe sequence.

Step 3: Work = Work + Allocation[P2] = (7 5 5) + (3 0 2) = (10 5 7)

.....P0.....P1.....P2.....P3.....P4.....

Finish = | true | true | true | true | true |

Step 4: Finish[Pi] = true for $0 \leq i \leq 4$

Hence, the system is in a safe state.

The safe sequence is <P1, P3, P4, P0, P2>.

Since the system is in safe state, the request can be granted.


Deadlock Detection

- ✓ If a system does not use either deadlock-prevention or deadlock-avoidance algorithm then a deadlock may occur.
- ✓ In this environment, the system must provide
 - An algorithm to examine the system-state to determine whether a deadlock has occurred.
 - An algorithm to recover from the deadlock.


Single Instance of Each Resource Type

- ✓ If all the resources have only a single instance, then deadlock detection-algorithm can be defined using a wait-for-graph.
- ✓ The wait-for-graph is applicable to only a single instance of a resource type.
- ✓ A wait-for-graph (WAG) is a variation of the resource-allocation-graph.
- ✓ The wait-for-graph can be obtained from the resource-allocation-graph by
 - removing the resource nodes and
 - collapsing the appropriate edges.
- ✓ An edge from P_i to P_j implies that process P_i is waiting for process P_j to release a resource that P_i needs.
- ✓ An edge $P_i \rightarrow P_j$ exists if and only if the corresponding graph contains two edges $P_i \rightarrow R_q$ and $R_q \rightarrow P_j$.

✓ For example:
Consider resource-allocation-graph and Corresponding wait-for-graph shown in below figure,



Resource-allocation-graph



Corresponding wait-for-graph.

- ✓ A deadlock exists in the system if and only if the wait-for-graph contains a cycle.
- ✓ To detect deadlocks, the system needs to
 - maintain the wait-for-graph and
 - periodically execute an algorithm that searches for a cycle in the graph.

Several Instances of a Resource Type

- ✓ The wait-for-graph is applicable to only a single instance of a resource type.
- ✓ However, the wait-for-graph is not applicable to a multiple instance of a resource type.
- ✓ The following detection-algorithm can be used for a multiple instance of a resource type.
- ✓ Assumptions:

Let 'n' be the number of processes in the system

Let 'm' be the number of resources types.

- ✓ Following data structures are used to implement this algorithm.

1. Available [m]

This vector indicates the no. of available resources of each type.

If Available[j]=k, then k instances of resource type R_j is available.

2. Allocation [n][m]

This matrix indicates no. of resources currently allocated to each process.

If Allocation[i,j]=k, then P_i is currently allocated k instances of R_j.

3. Request [n][m]

This matrix indicates the current request of each process.

If Request [i, j] = k, then process P_i is requesting k more instances of resource type R_j.

Step 1:

Let Work and Finish be vectors of length m and n respectively.

```

a) Initialize Work = Available
b) For i=0,1,2.....n
    if Allocation(i) != 0 then
        Finish[i] = false;
    else
        Finish[i] = true;

```

Step 2:

Find an index(i) such that both

```

a) Finish[i] = false
b) Request(i) <= Work.
If no such i exist, goto step 4.

```

Step 3:

```

Set:
Work = Work + Allocation(i)
Finish[i] = true

```

Go to step 2.

Step 4:

If Finish[i] = false for some i where $0 < i < n$, then the system is in a deadlock state.

Detection-Algorithm Usage

- ✓ The detection-algorithm must be executed based on following factors:
 1. The frequency of occurrence of a deadlock.
 2. The no. of processes affected by the deadlock.
- ✓ If deadlocks occur frequently, then the detection-algorithm should be executed frequently.
- ✓ Resources allocated to deadlocked-processes will be idle until the deadlock is broken.
- ✓ Deadlock occurs only when some processes make a request that cannot be granted immediately.
- ✓ The deadlock-algorithm must be executed whenever a request for allocation cannot be granted immediately.

- ✓ In this case, we can identify set of deadlocked-processes and specific process causing the deadlock.
 - ✓ The deadlock-algorithm must be executed in periodic intervals.
For example: once in an hour, whenever CPU utilization drops below certain threshold
-

Recovery from deadlock

Three approaches to recovery from deadlock:

- Inform the system-operator for manual intervention.
- Terminate one or more deadlocked-processes.
- Preempt(or Block) some resources.

Process Termination

- ✓ There are two methods to remove deadlocks:
 - Terminate all deadlocked-processes:
This method will definitely break the deadlock-cycle. However, this method incurs great expense. This is because, Deadlocked-processes might have computed for a long time. Results of these partial computations must be discarded. Probably, the results must be re-computed later.
 - Terminate one process at a time until the deadlock-cycle is eliminated:
This method incurs large overhead. This is because after each process is aborted, deadlock-algorithm must be executed to determine if any other process is still deadlocked
- ✓ For process termination, following factors need to be considered:
 - The priority of process.
 - The time taken by the process for computation & the required time for complete execution.
 - The no. of resources used by the process.
 - The no. of extra resources required by the process for complete execution.
 - The no. of processes that need to be terminated for deadlock-free execution.
 - The process is interactive or batch.

Resource Preemption

- ✓ Some resources are taken from one or more deadlocked-processes.
- ✓ These resources are given to other processes until the deadlock-cycle is broken.
- ✓ Three issues need to be considered:
 - **Selecting a victim**
 - Which resources/processes are to be pre-empted (or blocked)?
 - The order of pre-emption must be determined to minimize cost.
 - Cost factors includes
 - The time taken by deadlocked-process for computation.
 - The no. of resources used by deadlocked-process.
 - **Rollback**
 - If a resource is taken from a process, the process cannot continue its normal execution.
 - In this case, the process must be rolled-back to break the deadlock.
 - This method requires the system to keep more info. about the state of all running processes.
 - **Starvation**
 - In a system where victim-selection is based on cost-factors, the same process

may be always picked as a victim.

- As a result, this process never completes its designated task.
 - Ensure a process is picked as a victim only a (small) finite number of times.
-